

Virtualization with Containers and Kubernetes

Introduction with a focus on Docker

Andreas Roth

Training

09.06.2021

Contents

1 Virtualization

2 Containers

3 Docker

4 Kubernetes


Virtualization

Virtualization

Classical view of **a computer** running **an operating system**, running **a program**.

- OS is in full control of the hardware
- A program runs with lower privilege
- OS assigns resources to program

OS acts as a **supervisor** for programs.




Virtualization

Virtual Machine Monitor¹ (Hypervisor)

- 1 provides environment essentially identical to original machine
- 2 in full control of system resources
- 3 Programs run with minimal speed decreases in this environment

Virtual Machine¹


The environment for **programs** to run in, with a **Hypervisor** present, is called a **Virtual Machine**.



¹Popek G.J., Goldberg R. P. *Formal Requirements for Virtualizable Third Generation Architectures*, Communications of the ACM, Volume 7 (17), 1974


Containers

"Containerization" - A kind of virtualization...




OS-level virtualization¹

OS allows multiple instances of isolated user spaces (**containers**) to run applications in. They all share the OS Kernel.



- Processes inside the container cannot see anything outside (processes, files, ...). Physically, they are processes running on the same OS Kernel, however.
- There is no need for a virtualization layer below the Kernel to run containers, although there can be one!
- A **container engine** or **runtime** manages the containers

¹ OS-level Virtualization, Wikipedia 2021, https://en.wikipedia.org/wiki/OS-level_virtualization



Why?

- Bundle up an application **with all its dependencies!**
- More **lightweight and portable** than a virtual machine

How?

- Containers can be shipped as images (distribution infrastructure depends on system / engine)
- Who runs a container based on my image can be sure to have a **functionally identical environment**, independent of application type

How to make containers?

Advanced use of Linux features!

namespaces

- Limit what a process can **see**
- Change filesystem root with `chroot`
- Change apparent `pid`s
- ...

cgroups

- Limit what a process can **use**
- Hierarchical definition tree to assign quantities of memory, cpu cores, ... to processes and groups of processes

Definitely watch those talks:

- Containers from scratch (Liz Rice, 2018): <https://www.youtube.com/watch?v=8fi7uSYlOdc>
- Namespaces, cgroups and beyond: What are containers made from? (Jérôme Petazzoni, 2015): <https://www.youtube.com/watch?v=sK5i-N34im8>

There are / might be several software layers between the user interacting with a **container engine** and the **runtime** that actually runs containers.

Based on Linux namespaces and cgroups:

- [LXC](#)
- [Docker](#) (Docker engine, containerd)
- [Podman](#)

Other mechanisms / systems:

- [FreeBSD Jail](#) (Unix)
- [Solaris Zones](#) (Unix)
- [Windows Containers](#)

Docker

- Runtime and engine
- Building / managing of images / containers
- Distribution of images (Docker Hub)
- Interface to Windows and MacOS, different mechanisms
- (Orchestration of containers) swarm,
`docker-compose`
- Docker engine exists as open source, Docker Desktop and Docker Hub are developed by [Docker Inc.](#)

Docker is the most relevant container tool for us.

We will take a short tour of the most important **Docker concepts**:

- Running containers
- Building images with `Dockerfile`s
- Docker registries ([Docker Hub](#), company self-hosting solutions)
- Pushing / pulling images
- Docker volumes



- Feel free to try out the commands presented on the next slides during the training!
- You need a [Docker Desktop](#) installation (or [Docker Engine on Linux](#))
- As a preparation, clone github.com/scherbertlemon/docker-training

Container registries

... where prepared container images come from, if you do not build them based on your host system. They are usually organised in versioned image repositories.

[Docker Hub](#) hosted by Docker Inc.

- Open registry where you can store your images in **repositories**
- Public repositories are free of charge
- Many [officially curated images](#)

[Your company solution here](#)

- you might want to host your own registry

- Connect to a registry with `docker login <url>`. If no URL, Docker Hub is the default.
- Generate an **Access Token** if possible and use it for login instead of your password.

Images and Containers



Image

Read only content of a container, e.g. file system, environment variables, metadata

- Organized in **layers** ("differences")
- We usually **pull / push** images from / to **container registries**



Container

Extracted image content plus writable layer, that can be run on the host OS.

- We usually **run** processes in a container
- Changes to the file system exist as long as the container is not stopped and removed

Images and Containers



```
# get the official ubuntu image
docker pull ubuntu:focal
# check the local images
docker image ls
# check the properties of this image
docker image inspect ubuntu:focal
# remove the local image
docker image rm ubuntu:focal
```



```
# get an interactive shell in ubuntu
docker run -it --name ubu ubuntu:focal
# check the list of containers
docker container ls --all
# check the properties of ubu container
docker container inspect ubu
# remove stopped container
docker container rm ubu
```


Building images with Dockerfiles

```
FROM ubuntu:focal

LABEL maintainer="someone"

ENV TZ=Europe/Berlin

RUN apt update \
    && apt install \
    tzdata \
    --yes --no-install-recommends \
    && rm -rf /var/lib/apt/lists/*

WORKDIR /wdir
# goes to wdir
COPY localfile.txt .

ENTRYPOINT [ "date" ]
```

Example that prints the time when run

Dockerfile

... a recipe to perform certain tasks starting from an image, controlled by **directives**, to create a new, modified image.

- find this Dockerfile in `dockerfiles/time`
- [Dockerfile reference](#)
- [Best practises](#)
- Directives: `FROM`, `ENV`, `RUN`, `COPY`, `ENTRYPOINT`, `CMD`

Building images with Dockerfiles

```
FROM ubuntu:focal

LABEL maintainer="Someone"

ENV TZ=Europe/Berlin

RUN apt update \
    && apt install \
    tzdata \
    --yes --no-install-recommends \
    && rm -rf /var/lib/apt/lists/*

WORKDIR /wdir
# goes to wdir
COPY localfile.txt .

ENTRYPOINT [ "date" ]
```

Example that prints the time when run

```
# build an image from Dockerfile in current
# folder
$ docker build -t time:0.1 .
# run it as a container
$ docker run --rm time:0.1
Sun May 30 19:24:10 CEST 2021
# give an argument to date, UTC time
$ docker run --rm time:0.1 -u
Sun May 30 17:24:50 UTC 2021
```

- tag/name the image with `name:version`
- `.` indicates the **build context** (current folder)
- `--rm` deletes the container after it has stopped running

Building a more complex Dockerfile

Simple greeting app

- We install [miniconda](#) into our container.
 - We create a Python environment with `requirements.txt` in `/pysource/env`
 - We run the `flask` app defined in `greeting.py` in that environment
-
- navigate to `dockerfiles/greeting/0.1`
 - look at the Dockerfile
 - build&run it with
- ```
docker build -t greeting:0.1 .
docker run --rm -it -p 5000:5000 greeting:0.1
```
- `-p host:container` publishes ports from container to host.
  - Storage is just **in memory**:  
when app is restarted, counter is reset
- 
- Open your browser at [localhost:5000/hello/yourname](http://localhost:5000/hello/yourname)

# Running multiple containers

## Greeting app with database

- We need to run 2 containers: 1 for the app, 1 for the database
- We put them in a docker network `dbtest` and hook up the app to host port 5000 again

- navigate to `dockerfiles/greeting/0.2`
- different code than `greeting:0.1`
- build it with

```
docker build -t greeting:0.2 .
```

- Did this take as long as building `greeting:0.1`? → [build cache](#)
- Expects a postgres database on host `post:5432` with password `holymoly`

```
create a network
docker network create dbtest
run postgres container
docker run --rm --name post \
 --network dbtest -d \
 -e POSTGRES_PASSWORD=holymoly \
 postgres:latest
run app container
docker run --rm --name greet \
 --network dbtest -d -e PG_HOST=post \
 -p 5000:5000 greeting:0.2
check the network
docker inspect dbtest
```

# Docker volumes - persisting data

If we restart the `post` container, the recorded data is lost!

## Volumes

**Persisting data** by mounting persistent storage into the container file system.

Persistent storage can be

- **Bindings** to host file system
- **Named volumes**
- You can specify bind mounts similarly with `-v hostpath:containerpath`
- Also works for Windows paths, but performance may be inferior

Use a named volume for `post` container:

```
stop the container if still running
docker container stop post
re-run with named volume
docker run --rm --name post \
 --network dbtest -d \
 -e POSTGRES_PASSWORD=holyoly \
 -v pgdata:/var/lib/postgresql/data
 postgres:latest
look at what you did
docker volume ls
docker inspect pgdata
```

# Docker compose - running multi-container apps

- In order to tidy up our mess from before:

```
docker container stop greet post
docker network rm dbtest
```

- Is this not tedious? → `docker compose` !
- Navigate to `dockerfiles/greeting`
- Run to the same effect

```
run containers in background
leave -d for interactive
docker compose up -d
tidy up
docker compose down
```

- `docker-compose.yml` tells docker compose what to do

`dockerfiles/greeting/docker-compose.yml`

```
services:
 greet:
 image: greeting:0.2
 build: ./0.2/
 ports:
 - "5000:5000"
 environment:
 - PG_HOST=post
 post:
 image: postgres:latest
 volumes:
 - pgdata:/var/lib/postgresql/data
 environment:
 - POSTGRES_PASSWORD=holyoly
volumes:
 pgdata:
 external: True
```

# Pushing images to registry

- Assume you have built `greeting:0.2` locally.
- Create a repository on Docker Hub or your own registry: `yourname/greeting`
- Tag the image to push with that name and push it:

```
docker tag greeting:0.2 yourname/greeting:0.2
docker push yourname/greeting:0.2
```

- You need to be logged in to your registry (`docker login`)!

# Golden rules / best practises for containers

- **Keep images small!**  
(no unnecessary image content, delete setup files, [Multi-stage builds](#))
- **Every container should do one job, and do that well!**  
(no need to run multiple applications in the same container, e.g. webserver and database)
- **Don't use the writable layer of a container as storage, use volumes!**  
(if you store something in your container, bind the location either to a named volume or the host file system)
- **Use official images whenever you can!**  
([Official Images on Docker Hub](#). Inside a company network, you can at least use the Dockerfiles as guidance)
- **Adhere to the best practises for Dockerfiles!**  
[as outlined here](#)



# Docker tips for salvaging disk space

- Use the `prune` commands when you run out of disk space

```
with option --all also named local images will be deleted
docker image prune
in case you did not use --rm with docker run
docker container prune
or everything at once, except volumes
docker system prune
```

- **On Windows:** Pruning alone does not help, as the virtual disks of the Linux VM do not shrink.  
(Use Docker Desktop → Troubleshooting → Purge Data)
- **On Windows:** Use WSL 2 Backend, if possible!  
(Good look behind company proxy server, though!)




# Kubernetes

# Scaling applications / services


## Monolithic Architecture

*"Our applications run on one server (e.g. webserver and database). If you need to update, you have to power the whole system down."*



## Containerized architecture

*"We test our applications and ship them to our production server(s) with all dependencies. On update, downtime is really small."*



## Distributed, managed architecture

*"Our applications consist of independent microservices that communicate via interfaces, they can be replaced / updated without downtime. The system itself dynamically adapts to workload by scaling up/down."*






# Kubernetes

Kubernetes (K8s) according to the [documentation](#):



“Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation.”

kubectl



# Kubernetes objects

Just a few should be mentioned here:

- **Deployment**: packaged application consisting of several containers run in a **Pod**. Can be replicated for load balancing.
- **Service**: A connection to an application, "pod-agnostic" for availability
- **Pods**, **Nodes**, etc. are Kubernetes objects too! Get information about them with

```
info for all objects of one type
kubectl get object
kubectl describe object
info for a specific object
kubectl get object objectname
kubectl describe object objectname
```

- **Control Plane** controls a cluster of **Nodes**
- containers are deployed in **Pods** on **Nodes**

# Interacting with the cluster: kubectl

## kubectl

is a CLI tool to send commands to the **API server** for creating Kubernetes objects.

`kubectl` creates objects

- **imperatively**: Everything explicitly on the command line
- **declaratively**: with a **spec** representing the desired state (`.yaml` -file)

Let the cluster work to realize your spec:

```
kubectl apply -f kubernetes/greeting-depl-0-1.yaml
kubectl apply -f kubernetes/greeting-depl-0-2.yaml
```

Declarative way is preferred and more practicable!

# Create a greeting app deployment and service

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: greeting-depl
spec:
 selector:
 matchLabels:
 app: greetingapp
 replicas: 2 # 2 pods of template
 template:
 metadata:
 labels:
 app: greetingapp
 spec:
 containers:
 - name: greet
 image: greeting:0.1
 ports:
 - containerPort: 5000
```

**Example:** `kubernetes/greeting-depl-0-1.yaml`

```
kubectl apply \
 -f greeting-depl-0-1.yaml
see what just happened
kubectl get deployments
kubectl get pods
kubectl logs podname containername
```

Make available by creating a **service**:

```
kubectl expose deployment \
 greeting-depl \
 --type=NodePort \
 --port 5000
note the port it gets mapped to
kubectl get services
run proxy for access on localhost
kubectl proxy
```

# Update greeting app deployment

**Example:** `kubernetes/greeting-depl-0-2.yaml`

- Updated greeting image version
- Added postgres container

Declare the new spec:

```
kubectl apply -f greeting-depl-0-2.yaml
```

- Pods are replaced one by one, without the service terminating!
- Eventually, you will see the newer greeting app.

Tidy up:

```
kubectl delete service greeting-depl
kubectl delete deployment greeting-depl
```

Containers and pods are terminated!



# Concluding notes

- replicas > 1: we do not know which of our apps we will get (load balancing, no persistent storage)
- persistent storage is possible, here only stateless applications were shown
- **Not mentioned here:** security, encryption, secrets
- Docker Desktop users can activate a one-node kubernetes cluster in the settings
- Another one-node try-out solution: [Minikube](#)

## Further reading:

- [Kubernetes documentation](#): Includes lots of learning resources and examples.
- [Play with Kubernetes classroom](#): Interactive training session giving an overview over concepts.

Thank you for participating!